

Pipelining

Nachiket Kapre

nachiket@uwaterloo.ca



Outline

- ▶ Need for pipelining
 - ▶ CPU model (IPC, CPI)
- ▶ Latency and Throughput
- ▶ Pipelining Syntax
- ▶ Bubbles in a pipeline

Idea of pipelining

- ▶ Modern automated factories have *assembly lines*
- ▶ Each factory worker/robot specializes in one aspect of the manufacturing flow
- ▶ Allows us to design **high-throughput** pipelines of product manufacturing
- ▶ Compare: skilled craftsman who delivers entire product

Model T assembly



Electronics assembly



https://commons.wikimedia.org/wiki/File:Jinbei_production_line.jpg

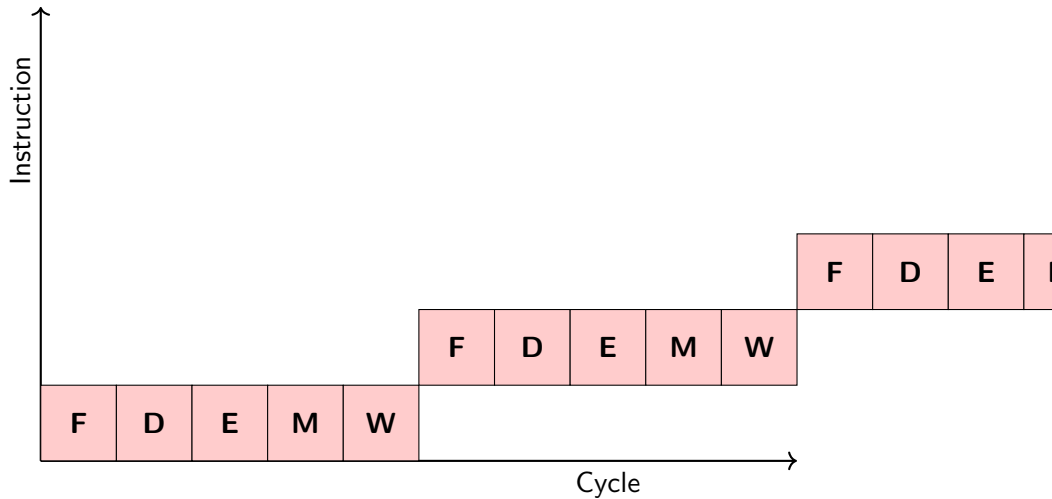
Pipelining in Modern CPUs

- ▶ C code compiled into a **sequence** of instructions (data movement + arithmetic)
- ▶ Simple CPU pipelines can pack entire instruction processing in a single cycle
 - ▶ **Fetch** → **Decode** → **Execute** → **Memory** → **Writeback**
- ▶ Better CPUs pipeline each stage into a cycle.
 - ▶ Pipelining allows **overlapping** of computation for different instructions in the same clock cycle
 - ▶ **But**, must check for dependencies/hazards.

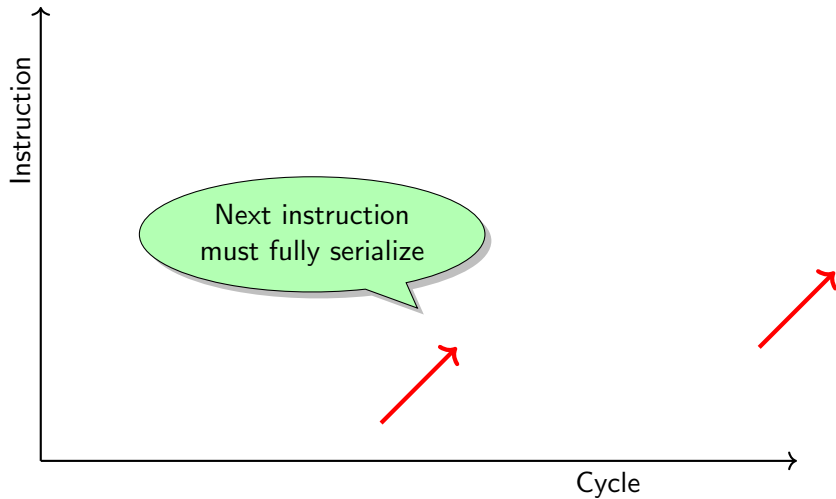


Fetch → Decode → Execute → Memory → Write


Execution Flow on a Unpipelined CPU



Execution Flow on a Unpipelined CPU



Execution Flow on a Pipelined CPU (no dependencies)



Advantages of Pipelining a CPU


- ▶ Faster clock frequency → fewer logic gate per stage
- ▶ Improved throughput of instructions retired
- ▶ **IPC** Instruction Per Cycle is ↑, **CPI** Clock cycles per instruction ↓
- ▶ Not a free lunch → pay extra in silicon area + power
- ▶ Also, when dependencies occur, you have to **stall** the pipeline
 - ▶ Order of instruction execution is important
 - ▶ Previous instruction must finish before next **dependent** instruction is issued

Execution Flow on a Pipelined CPU (poly with dependencies)

poly:

```
imull    %a, %x, %t1
imull    %b, %x, %t2
imull    %t1, %x, %t3
addl     %t2, %t3, %t4
addl     %c, %t4, %t5
movl     %t5, %y
ret
```

Execution Flow on a Pipelined CPU (with dependencies)



Understanding Pipelining in CPU

- ▶ CPU operation is internally pipelined (often, quite extensively)
- ▶ Dependencies cause **pipeline stalls** (or **bubbles**) in the pipeline
- ▶ **CPU has special hardware to track dependencies** →
 - ▶ In-order processing: Each instruction checks if input operands used in previous k instruction window. k depends on how depth of CPU pipeline.
 - ▶ Out-of-order processing: Tomasulo's algorithm used for out-of-order processors where multiple instructions M can be issued together. Must track dependencies across $k*M$ possible set of instructions.
- ▶ When designing your own custom hardware, you can reason about dependencies yourself!


Pipelining Terminology

- ▶ When measuring CPU performance, we count
 - ▶ (1) *INST* total number of instructions retired/executed
 - ▶ (2) *CYC* total number of cycles required
- ▶ Metric of goodness
 - ▶ **IPC** Instructions Per Cycle = $\frac{INST}{CYC}$
 - ▶ Another metric is **CPI** Clock Cycles Per Instruction = $\frac{CYC}{INST}$
 - ▶ Thus, **IPC** = $1/\text{CPI}$
- ▶ For unpipelined CPUs, **CPI** is depth of pipeline.
 - ▶ Here, the metric value does not change with dependencies
 - ▶ Paying the worst-case waiting time for each instruction
- ▶ For pipelined single-issue CPUs, **IPC** is 1 (best case).
 - ▶ Application dependencies will affect the metric in this case
 - ▶ If no dependency, IPC of 1 is possible. With dependencies, IPC will be < 1

Analyzing CPU performance (Unpipelined CPU)

- ▶ Observed counts
 - ▶ **CYC** = $5 * k$ cycles
 - ▶ **INST** = k instructions
- ▶ Computation of performance metrics
 - ▶ **IPC** = $\frac{k}{5 * k} = \frac{1}{5}$
 - ▶ **CPI** = $\frac{5 * k}{k} = 5$
- ▶ Design is slow, but predictable

Execution Flow on a Pipelined CPU (with dependencies)



Analyzing CPU performance

- ▶ Observed counts
 - ▶ **CYC** = 14 cycles
 - ▶ **INST** = 5 instructions
- ▶ Computation of performance metrics
 - ▶ **IPC** = $\frac{5}{14}$
 - ▶ **CPI** = $\frac{14}{5}$
- ▶ Design is faster, but unpredictable (each application may run differently as per dependencies)

Effect of Pipelining

- ▶ **Deeper** pipelines → lot of register stages from input to output
- ▶ More registers → High Latency
- ▶ More registers → High Throughput
- ▶ More registers → Smaller clock period
- ▶ More registers → More area

Effect of Pipelining


- ▶ Beyond a certain number of registers, clock period does not improve much → no more throughput improvements
- ▶ Only impact is higher latency and more area.
- ▶ **Challenge:** Find the sweet spot

Pipelining in poly

- ▶ Manually decompose the $a \times x^2 + b \times x + c$ into smaller operations
- ▶ 1–4 pipeline stages needed to hit granularity of individual operation
- ▶ Frequency is determined by the length of logic delays from input register to output register
- ▶ Note the delayed inputs – assume all inputs arrive at the same time. Important to delay signals to align their arrival at operator.

Pipelining Notation

- ▶ Show inputs and outputs with boxes, show arithmetic and logic operations in circles, show muxes with standard notation
- ▶ Pipelining stages are indicated with vertical dotted lines
 - ▶ If it helps you can label each stage by count (pipeline stage 1,2,3. . .)
- ▶ Remember, circuit edges that cross vertical pipeline edges **must** have a register for **each** cut.




Pipelining Terminology

- ▶ **Throughput:** Rate at which circuit can consume inputs.
 - ▶ The highest throughput possible is 1, where you can push new inputs into the circuit every clock cycle.
 - ▶ If you can consume inputs once every two clocks, the throughput is $\frac{1}{2}$.
 - ▶ Throughput is not defined for combinational circuits. Need a periodic clock signal to define throughput.
- ▶ **Latency:** Time required for the first output to emerge from the circuit for a given input.
 - ▶ This is the end-to-end latency for one item of data to propagate from inputs of the circuit to its output.
 - ▶ Ideally, you want low latency, a value of 1 is minimum for a pipelined circuit.
 - ▶ A purely combinational output has a latency of 0
- ▶ **Clock Period:** Time separation between clock edges.
 - ▶ A faster frequency will results in small clock periods.
 - ▶ Ideally, you want to run your circuit as fast as required by the design specification.

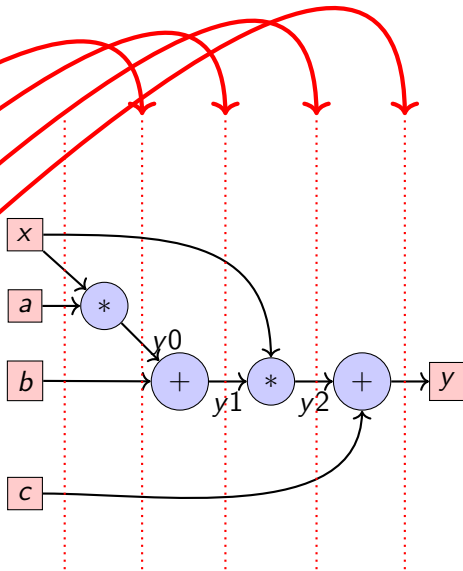
Pipelining poly circuit view

```
always @(posedge clk) begin
  if(rst) begin
    y0 <= {16{1'b0}};
    y1 <= {16{1'b0}};
    y2 <= {24{1'b0}};
    y <= {24{1'b0}};
    x_r1 <= {8{1'b0}};
    x_r2 <= {8{1'b0}};
    b_r1 <= {8{1'b0}};
    c_r1 <= {8{1'b0}};
    c_r2 <= {8{1'b0}};
    c_r3 <= {8{1'b0}};
  end else begin
    // stage 1
    y0 <= a_r * x_r;
    x_r1 <= x_r;
    b_r1 <= b_r;
    c_r1 <= c_r;
    // stage 2
    y1 <= y0 + b_r1;
    x_r2 <= x_r1;
    c_r2 <= c_r1;
    // stage 3
    y2 <= y1 * x_r2;
    c_r3 <= c_r2;
    // stage 4
    y <= y2 + c_r3;
  end
end
```




Pipelining poly circuit view

```
always @(posedge clk) begin
  if(rst) begin
    y0 <= {16{1'b0}};
    y1 <= {16{1'b0}};
    y2 <= {24{1'b0}};
    y <= {24{1'b0}};
    x_r1 <= {8{1'b0}};
    x_r2 <= {8{1'b0}};
    b_r1 <= {8{1'b0}};
    c_r1 <= {8{1'b0}};
    c_r2 <= {8{1'b0}};
    c_r3 <= {8{1'b0}};
  end else begin
    // stage 1
    y0 <= a_r * x_r;
    x_r1 <= x_r;
    b_r1 <= b_r;
    c_r1 <= c_r;
    // stage 2
    y1 <= y0 + b_r1;
    x_r2 <= x_r1;
    c_r2 <= c_r1;
    // stage 3
    y2 <= y1 * x_r2;
    c_r3 <= c_r2;
    // stage 4
    y <= y2 + c_r3;
  end
end
```




Pipelining poly circuit view

```
always @(posedge clk) begin
    if(rst) begin
        y1 <= {16{1'b0}};
        y <= {24{1'b0}};
        x_r1 <= {8{1'b0}};
        c_r1 <= {8{1'b0}};
    end else begin
        // stage 1
        y1 <= a_r * x_r + b_r;
        x_r1 <= x_r;
        c_r1 <= c_r;
        // stage 2
        y <= y1 * x_r1 + c_r1;
    end
end
```




Pipelining poly circuit view

```
always @(posedge clk) begin
  if(rst) begin
    y1 <= {16{1'b0}};
    y <= {24{1'b0}};
    x_r1 <= {8{1'b0}};
    c_r1 <= {8{1'b0}};
  end else begin
    // stage 1
    y1 <= a_r * x_r + b_r;
    x_r1 <= x_r;
    c_r1 <= c_r;
    // stage 2
    y <= y1 * x_r1 + c_r1;
  end
end
```




Pipelining poly circuit view

```
always @(posedge clk) begin
  if(rst) begin
    y <= {24{1'b0}};
  end else begin
    // stage 1
    y <= (a_r * x_r + b_r) * x_r + c_r;
  end
end
```



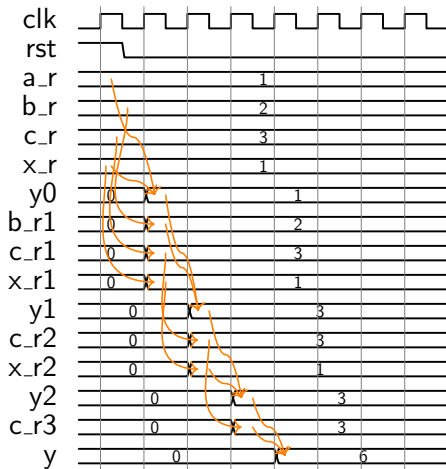
Pipelining poly circuit view

```
always @(posedge clk) begin
  if(rst) begin
    y <= {24{1'b0}};
  end else begin
    // stage 1
    y <= (a_r * x_r + b_r) * x_r + c_r;
  end
end
```




Pipelining Waveforms (4-Stage Pipeline)

```
always @(posedge clk) begin
  if(rst) begin
    y0 <= {16{1'b0}};
    y1 <= {16{1'b0}};
    y2 <= {24{1'b0}};
    y <= {24{1'b0}};
    x_r1 <= {8{1'b0}};
    x_r2 <= {8{1'b0}};
    b_r1 <= {8{1'b0}};
    c_r1 <= {8{1'b0}};
    c_r2 <= {8{1'b0}};
    c_r3 <= {8{1'b0}};
  end else begin
    // stage 1
    y0 <= a_r * x_r;
    x_r1 <= x_r;
    b_r1 <= b_r;
    c_r1 <= c_r;
    // stage 2
    y1 <= y0 + b_r1;
    x_r2 <= x_r1;
    c_r2 <= c_r1;
    // stage 3
    y2 <= y1 * x_r2;
    c_r3 <= c_r2;
    // stage 4
    y <= y2 + c_r3;
  end
end
```




Pipelining Waveforms (2-Stage Pipeline)

```
always @(posedge clk) begin
  if(rst) begin
    y1 <= {16{1'b0}};
    y <= {24{1'b0}};
    x_r1 <= {8{1'b0}};
    c_r1 <= {8{1'b0}};
  end else begin
    // stage 1
    y1 <= a_r * x_r + b_r;
    x_r1 <= x_r;
    c_r1 <= c_r;
    // stage 2
    y <= y1 * x_r1 + c_r1;
  end
end
```




Pipelining Waveforms (1-Stage Pipeline)

```
always @(posedge clk) begin
  if(rst) begin
    y <= {24{1'b0}};
  end else begin
    // stage 1
    y <= (a_r * x_r + b_r) * x_r + c_r;
  end
end
```




Pipelining poly analysis

- ▶ Latency = 5
- ▶ Throughput = 1
- ▶ Cost = $(8+8+8+8) + (16+8+8+8) + (16+8+8) + (24+8) + 24 = 160$
- ▶ Each circuit edge that crosses vertical line requires a FF/register
- ▶ Clock Period = $\max(T_{*1}, T_{+1}, T_{*2}, T_{+2})$
- ▶ General expression is max over all FF \rightarrow FF paths




Pipelining poly analysis

- ▶ Latency = 3
- ▶ Throughput = 1
- ▶ Cost = $(8+8+8+8) + (16+8+8) + 24 = 88$
- ▶ Clock Period = $\max(T_{*1} + T_{+1}, T_{*2} + T_{+2})$



Pipelining poly analysis


- ▶ Latency = 2
- ▶ Throughput = 1
- ▶ Cost = $(8+8+8+8) + 24 = 56$
- ▶ Clock Period =
 $T_{*1} + T_{+1} + T_{*2} + T_{+2}$
- ▶ **Note:** Throughput stayed 1 in all cases



Wire Pipelining


- ▶ On modern FPGAs and ASICs, most of the delay is in the wire, and not logic
- ▶ For instance, gates/LUTs may take 0.2–0.3 ns while a wire connecting the output to next gate may take 1–2 ns.
- ▶ Increasingly, we must extend the idea of pipelining to wires.

Wire pipelining poly circuit view



Clock Period =

Wire pipelining poly circuit view



Clock Period =


Wire Pipelining Code

```
// state 1
y0 <= a * x;
x_r <= x;
b_r <= b;
c_r <= c;
// stage 2
y1 <= y0 + b_r;
x_r1 <= x_r;
c_r1 <= c_r;
// wire delays
y1_r <= y1;
x_r2 <= x_r1;
c_r2 <= c_r1;
y1_r1 <= y1_r;
x_r3 <= x_r2;
c_r3 <= c_r2;
y1_r2 <= y1_r1;
x_r4 <= x_r3;
c_r4 <= c_r3;
// stage 3
y2 <= y1_r2 * x_r4;
c_r5 <= c_r4;
// stage 4
y <= y2 + c_r5;
```

- ▶ In Verilog, you need to add dummy pipeline stages → they do no compute, but provide freedom to the FPGA CAD tools to split wires
- ▶ Wirelength is only known after full place-and-route.
- ▶ Synthesis tools aim to *estimate* wire lengths and delays, but a few iterative compilations are required regardless.
 - ▶ Similar to profile-guided optimization in gcc
- ▶ **Impact:**
 - ▶ Must design **parametric** hardware
 - ▶ Retiming can automate this

Wire Pipelining Waveform

```
// state 1
y0 <= a * x;
x_r <= x;
b_r <= b;
c_r <= c;
// stage 2
y1 <= y0 + b_r;
x_r1 <= x_r;
c_r1 <= c_r;
// wire delays
y1_r <= y1;
x_r2 <= x_r1;
c_r2 <= c_r1;
y1_r1 <= y1_r;
x_r3 <= x_r2;
c_r3 <= c_r2;
y1_r2 <= y1_r1;
x_r4 <= x_r3;
c_r4 <= c_r3;
// stage 3
y2 <= y1_r2 * x_r4;
c_r5 <= c_r4;
// stage 4
y <= y2 + c_r5;
```



DRIVE stage in the Pentium-4 pipeline


Architectural Fixes: Pentium-4

1	
2	TC Next IP
3	
4	TC Fetch
5	Drive
6	Alloc
7	
8	Rename
9	Queue
10	Schedule 1
11	Schedule 2
12	Schedule 3
13	Dispatch 1
14	Dispatch 2
15	Register File 1
16	Register File 2
17	Execute
18	Flags
19	Branch Check
20	Drive

Pipeline stages
dedicated to driving
signals across chip

Image removed due to copyright restrictions.

Extensive wire pipelining in Intel Hyperflex FPGA



- Registers are available in every routing segment
- Registers are available on all block inputs (ALM, M20K blocks, DSP blocks, and I/O cells)

<https://www.intel.com/content/dam/www/programmable/us/en/pdfs/literature/wp/>

wp-01231-understanding-how-hyperflex-architecture-enables-high-performance-systems.pdf

Pipeline bubbles


- ▶ In water pipes, bubbles arise due to air intake → bubble smoothly rides with the flow and is released.
- ▶ In CPUs, processing pipeline **bubbles** arise due to stalls (which in turn are due to data dependencies).
- ▶ In custom datapaths, bubbles arise due to resource sharing (next lecture), or absence of input data.
 - ▶ Resource sharing is known upfront when designing the hardware
 - ▶ Data arrival uncertainty is now known at compile/design time
- ▶ **Key:** A bubble is a NULL operation through the hardware pipeline. Outputs of the pipeline stage when the bubble is passing through is ignored. Bubble moves through the pipeline in LATENCY cycles.

Implementing bubbles

- ▶ Bubbles are simply **invalid** data in a pipeline
- ▶ Can be trivially handled with a special **valid** indication
- ▶ For resource sharing , **valid** signal generated by state machine
- ▶ For absent data, **valid** signal must be dynamically determined
 - ▶ Tag input with valid, and output with valid
 - ▶ Generate a shift register with identical delay as pipeline latency
 - ▶ Feed input valid into the shift register, and connect output valid to output of shift register
- ▶ Must handle multiplexers (if present) carefully → need a separate shift register for each branch and select input, and combine the valids based on select appropriately


Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```




Implementing bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```



Waveform for bubbles

```
always @(posedge clk) begin: valid
  if(rst) begin
    v <= {3{1'b0}};
    y_v <= 1'b0;
  end else begin
    v[0] <= x_v; //
    v[1] <= v[0];
    v[2] <= v[1];
    v[3] <= v[2];
    y_v <= v[3]; //
  end
end
```



Wrapup

- ▶ Datapath pipelining useful to improve throughput at the expense of latency
- ▶ Unlike CPU pipelines, custom datapath pipelines allow a designer to choose throughput, latency, clock period combinations to meet system requirements
- ▶ Wire delays more important in modern chips → pipelining must consider wires.
- ▶ Bubbles in the input must be handled carefully to avoid corrupting output data
- ▶ **Question:** Can we pipeline any circuit?
 - ▶ Fundamental limit to how much pipelining is possible → single LUT
 - ▶ Feedforward circuits can always be pipelined