

# Multi-Layer Networks

Tripp Deep Learning F22



## **TODAY'S GOAL**

---

By the end of the class, you should be able to explain how multi-layer networks work and what they are capable of, and you should be able to choose suitable outputs and losses for regression and classification.


---

# Summary


1. Multi-layer feedforward networks make a series of simple but unintuitive calculations
2. Multi-layer networks are composite functions
3. Multi-layer networks break the curse of dimensionality
4. A two-layer network can closely approximate any continuous function on a bounded domain
5. Networks with different weight matrices can perform identical mappings
6. Classification and regression employ different output and loss functions
7. Multi-layer networks are also called multi-layer perceptrons, but are not perceptrons

**MULTI-LAYER FEEDFORWARD NETWORKS  
MAKE A SERIES OF SIMPLE BUT  
UNINTUITIVE CALCULATIONS**

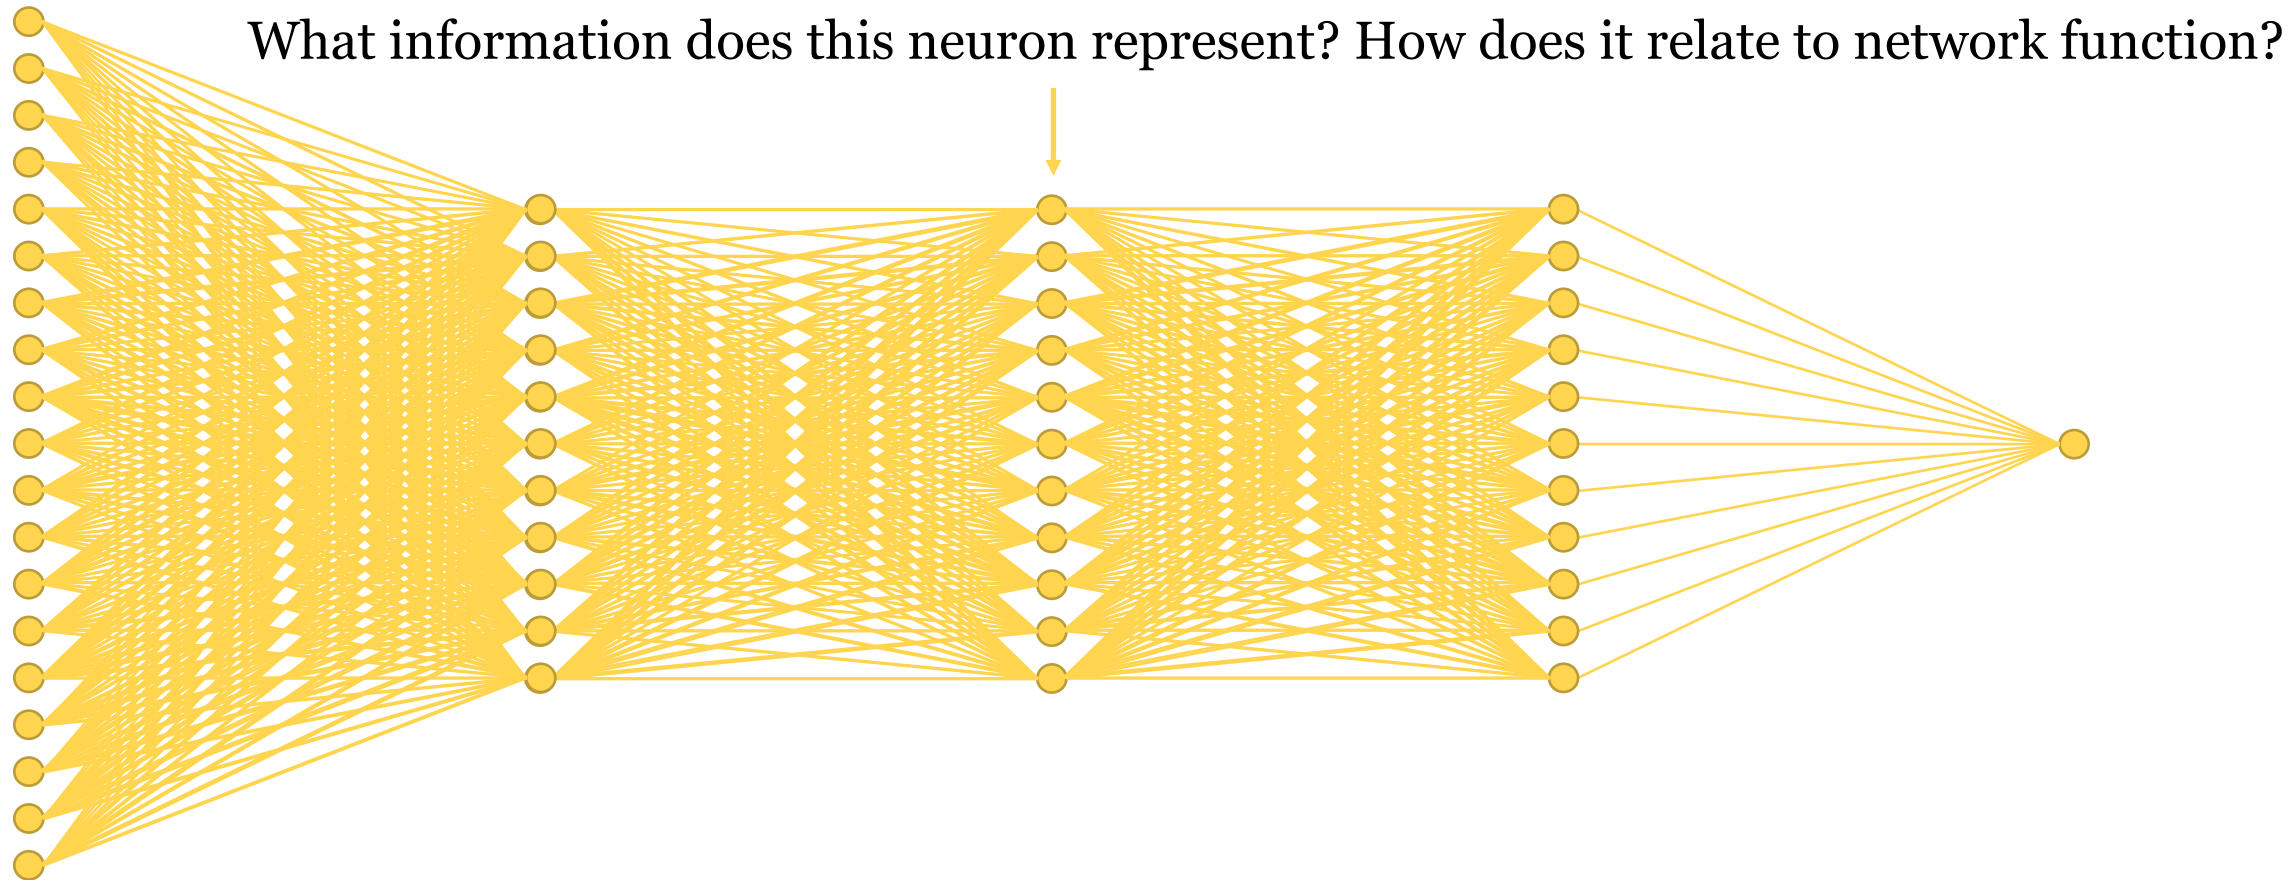
# Overall network function is intuitive



# How each neuron works is intuitive



# Role of each neuron in network function is opaque



# **MULTI-LAYER NETWORKS ARE COMPOSITE FUNCTIONS**




$x$


$f_1(x)$

$f_2(f_1(x))$

$f_3(f_2(f_1(x)))$

$f_4(f_3(f_2(f_1(x))))$






$$\begin{aligned}
 y &= g_y(w_1^y h_1 + w_2^y h_2 + b_y) \\
 &= g_y(w_1^y g_h(w_{11}^h x_1 + w_{12}^h x_2 + b_{h1}) + w_2^y g_h(w_{21}^h x_1 + w_{22}^h x_2 + b_{h2}) + b_y)
 \end{aligned}$$

# **MULTI-LAYER NETWORKS BREAK THE CURSE OF DIMENSIONALITY**

# Recall: Curse of dimensionality




- Multi-layer neural networks take advantage of this by adapting basis functions so that
  - Regions of variation correspond to regions over which input typically varies
  - Directions of variation correspond to directions over which output typically varies



# Basis functions adapt to the task



- Example:  
Approximating different functions with three sigmoid hidden units and a linear output unit
- The hidden units adapt to the task

$$x \quad h_i \quad \hat{y} = \sum w_i h_i + b$$



..... target  
—— approx

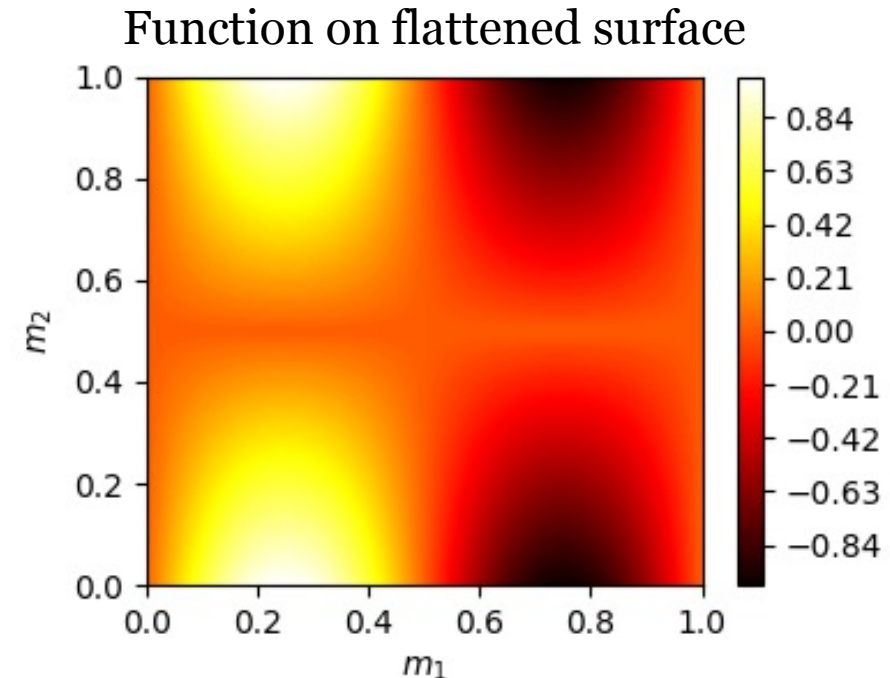
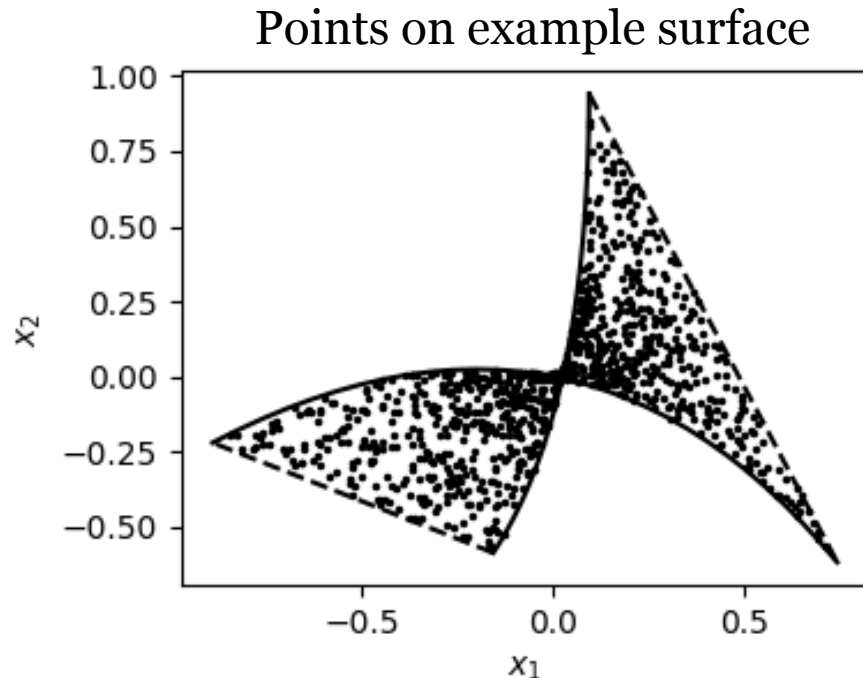
adaptive  
basis  
functions



Bishop, Pattern Recognition & Machine Learning

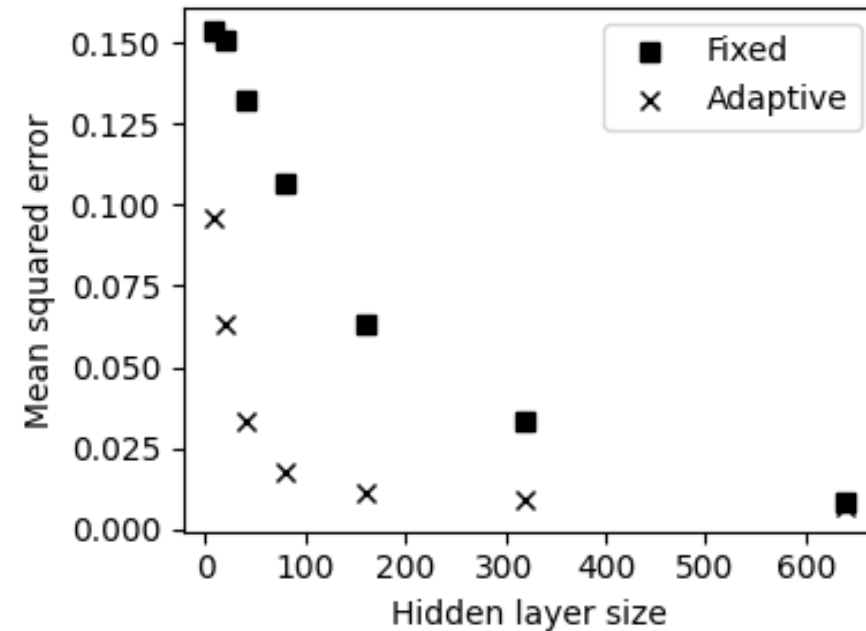

# Fewer basis functions are needed if they are adaptive

- Example: Nonlinear function on randomly bent 2D surface in 5D space
- Different parts of the surface can overlap in lower-dimensional projections, so all the dimensions should be considered



# Fewer basis functions are needed if they are adaptive

- Two-layer networks with  $n$  ReLU neurons in the hidden layer



**A TWO-LAYER NETWORK CAN CLOSELY  
APPROXIMATE ANY CONTINUOUS  
FUNCTION ON A BOUNDED DOMAIN**



# Universal approximation

- In principle a network with a single hidden layer can approximate any continuous function with arbitrary precision, provided it has enough hidden units
- Proofs allow for an unbounded number of hidden units (although there has also been separate work to establish bounds)
- The number of hidden units required may be impractical
- This result doesn't guarantee that a given algorithm can learn the approximation

# Universal approximation

Two-layer networks of the form,

$$\hat{y} = \mathbf{w}_y^T [g(\mathbf{w}_i^T \mathbf{x} + b_i)] + b_y$$

where  $g$  is a scalar nonlinearity that can be a sigmoid or ReLU (or many other continuous functions) can approximate any function over a bounded domain to arbitrary precision, i.e.,

$$|\hat{y} - y| < \varepsilon$$

for all  $\mathbf{x}$  in the domain and any choice of  $\varepsilon$ .

# Influence of sigmoid hidden units with one input


- Let's get a sense of how this can be
- Suppose we have a two-layer network  $G$  with one input  $x$ ,  $n$  sigmoid hidden units and one linear output,

$$\hat{y} = G(x)$$


$$= \sum_{i=1}^n w_{yi} (1 + e^{-(w_{ix}x + b_i)})^{-1} + b_y$$

- Then  $\hat{y}$  is a sum of functions that can vary in slope, offset, and height, as illustrated in these figures


$$(1 + e^{-x})^{-1}$$




$$(1 + e^{-w_{ix}x})^{-1}$$



$$(1 + e^{-(x+b_i)})^{-1}$$




$$w_{yi}(1 + e^{-x})^{-1}$$







# Approximating a 1D function in steps

- To approximate a function  $y = f(x)$  on the interval  $[x_a, x_b]$  so that  $|\hat{y} - y| \leq \varepsilon$ :
  1. Choose a step size  $\Delta x$  and choose a weight  $w_{ix}$  large enough that the sigmoid changes most of the way from 0 to 1 over a span of  $\Delta x$
  2. Set the bias to  $b_y = f(x_a)$
  3. For each  $x_i$  from  $x_a$  to  $x_b$  in steps of  $\Delta x$ :
    1. Add a hidden neuron with  $w_{yi} = f(x_i + \Delta x) - f(x_i)$  and  $b_i = -w_{ix}(x_i + \Delta x/2)$
- This will make  $\hat{y}$  approximate  $y$ , perhaps poorly. To make  $|\hat{y} - y| \leq \varepsilon$  increase  $w_{ix}$  and decrease  $\Delta x$  as needed



# Extending to multiple input dimensions

- Approximate the desired function with a multi-dimensional Fourier series
  - Use a finite but arbitrarily large number of frequency components
  - Use real-valued form (e.g., amplitude-phase)
- Approximate each component as before
  - Each component is a sinusoidal function of the projection onto a vector  $\mathbf{v}$  in the input space
  - Approximate each sinusoid with sigmoid functions that point along  $\mathbf{v}$ , specifically  $\left(1 + e^{-a\mathbf{v}^T \mathbf{x} + b}\right)^{-1}$



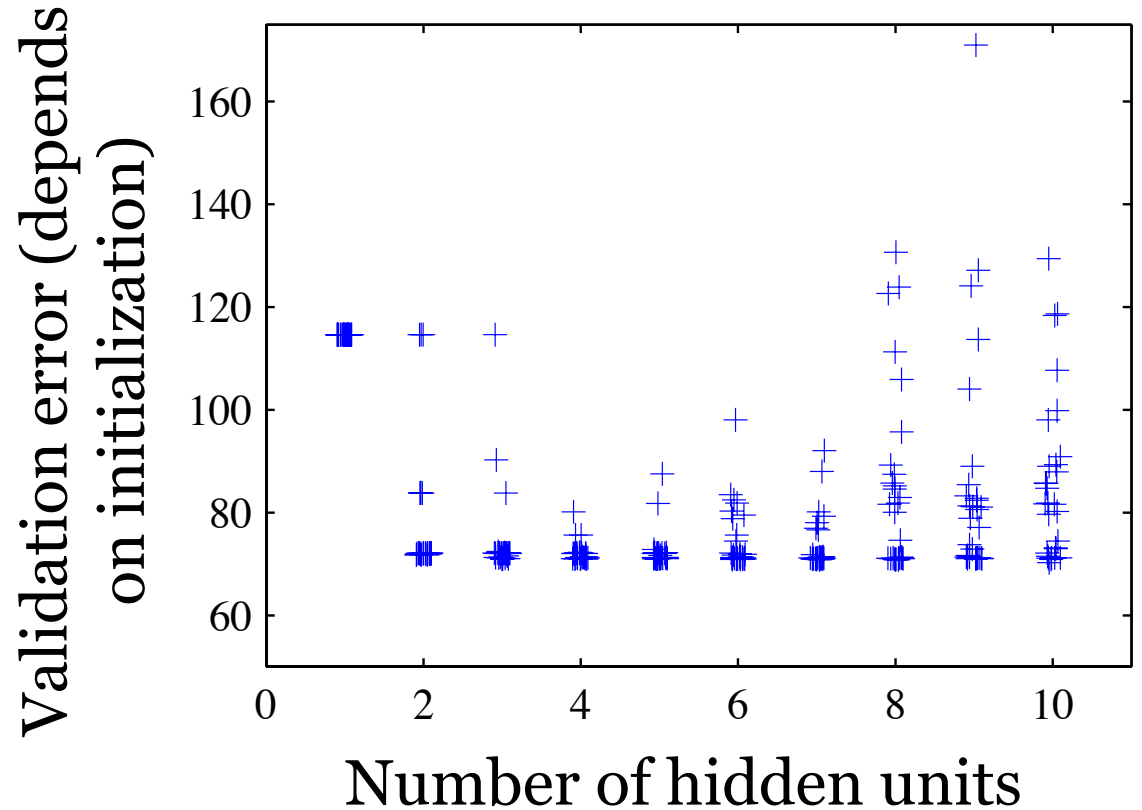
# Proofs



- For a discussion of various proofs: Scarselli, F., & Tsoi, A. C. (1998). Universal approximation using feedforward neural networks: A survey of some existing methods, and some new results. *Neural Networks*, 11(1), 15-37.

# Large hidden layer increases risk of overfitting


- These are results of training small networks with 30 sets of random initial weights for each hidden-layer size



Goodfellow, Bengio & Courville, *Deep Learning*

# Depth

- Deeper networks can represent some kinds of functions more efficiently, particularly functions that are compositions of simpler functions
- Choosing a deep architecture is consistent with the belief that the desired mapping is in this category
- Empirically, greater depth results in better generalization in a wide range of tasks



Goodfellow, Bengio  
& Courville, *Deep  
Learning* (pg. 197)



# **NETWORKS WITH DIFFERENT WEIGHT MATRICES CAN PERFORM IDENTICAL MAPPINGS**

# Invariance to parameter changes

- For any input-output function that a network may perform with a certain set of parameters, there are many other sets of parameters that implement *exactly* the same function
- In particular, swapping the input and output weights and the bias of a hidden unit with those of another hidden in the same layer, with the same nonlinearity, does not affect the output
- Some networks have additional invariances, for example
  - With tanh hidden-layer activation functions,  $\tanh(-a) = -\tanh(a)$ , so changing the signs of both the input and output weights together has no effect
  - With multi-head attention in transformers, the heads are interchangeable
  - With ReLU hidden-layer activation functions, some neurons can become “dead” (unresponsive to any input), so continuous changes in their parameters within some neighbourhood have no effect
  - With ReLU hidden-layer activation functions, multiplying all the weights and biases of one layer by  $\alpha$  and dividing the weights of the next layer by  $\alpha$  has no effect

# Invariance to parameter changes

- This means that there isn't a single global minimum of the loss
- For any minimum, there are **many** equivalent minima
  - E.g., for a network with a single hidden layer with ten units, at least  $10! = 3,628,800$  equally optimal minima

# **CLASSIFICATION AND REGRESSION EMPLOY DIFFERENT OUTPUT AND LOSS FUNCTIONS**



# Using negative log likelihood for loss

We want to perform maximum likelihood estimation of the network parameters,  $\theta$ . In other words, we adjust the parameters to maximize the probability of the target given the input and the parameters,

$$p(\mathbf{y}|\theta) = \prod_{n=1}^N p(y_n|\theta),$$

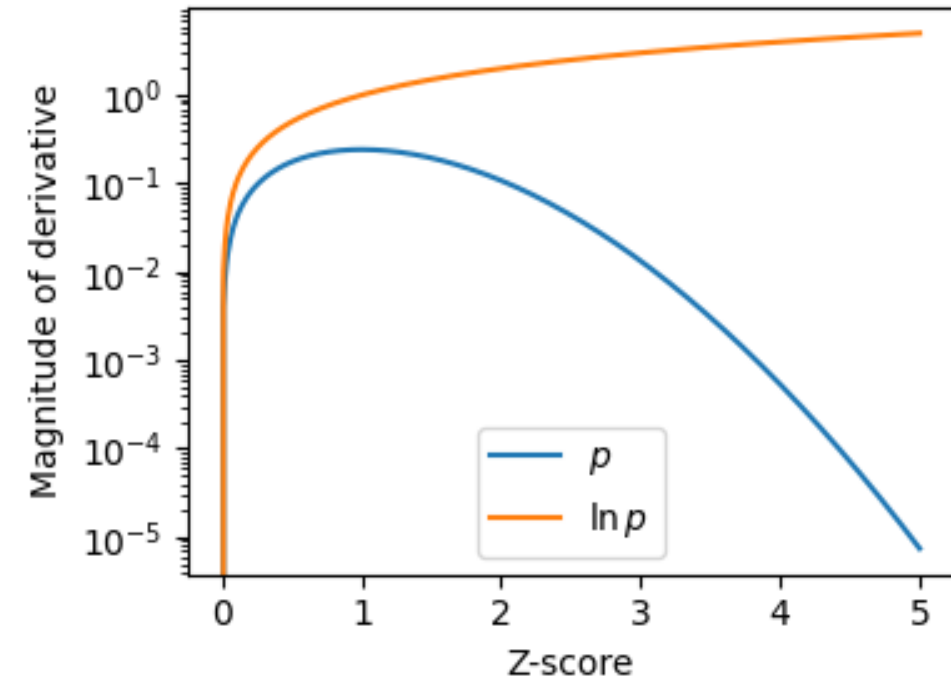
over  $N$  points of training data. This is equivalent to minimizing the negative log likelihood,

$$L(\theta) = -\ln p(\mathbf{y}|\theta) = -\sum_{n=1}^N \ln p(y_n|\theta).$$



# Using negative log likelihood for loss

- While they have the same minima, negative log probability is better than negative probability for several reasons:
  - It is a sum rather than a product, so we can do gradient descent on parts of it at a time (minibatches) and these actions combine properly
  - Multiplying multiple small probabilities can underflow the floating-point representation, whereas their log takes on a better-behaved range of values
  - It's useful to take gradient steps that scale with the gradient of the -ve log probability, whereas doing this with probabilities would result in very small steps far from the optimum (plot is for Gaussian distribution)



# Using negative log likelihood for loss

Example: In a binary classification network with one output, the target is  $y \in \{0,1\}$  and we want the network's output  $\hat{y}$  to be the probability that  $y = 1$ . Over  $N$  examples in the training dataset, the likelihood is,

$$p(\mathbf{y}|\boldsymbol{\theta}) = \prod_{n=1}^N \hat{y}_n^{y_n} (1 - \hat{y}_n)^{1-y_n},$$

The negative log-likelihood is,

$$L(\boldsymbol{\theta}) = -\ln p(\mathbf{y}|\boldsymbol{\theta}) = -\sum_{n=1}^N y_n \ln \hat{y}_n + (1 - y_n) \ln(1 - \hat{y}_n),$$

the binary cross-entropy loss that we have already seen.


# Derivative of the loss for binary classification

For a binary classification network with one output, the output for the  $n^{th}$  data point is,

$$\hat{y}_n = \sigma(a_n) = \frac{1}{1 + \exp(-a_n)},$$

where  $a_n$  is a linear combination of the activities of the second-last layer.

The derivative of this  $\sigma$  function is  $\sigma(1 - \sigma)$ .





# Derivative of the loss for binary classification

The loss for the  $n^{th}$  data point is,

$$L = -y_n \ln \hat{y}_n - (1 - y_n) \ln(1 - \hat{y}_n).$$

The derivative of the loss with respect to  $a_n$  is,

$$\begin{aligned} \frac{dL}{da_n} &= -\frac{y_n}{\hat{y}_n} \hat{y}_n(1 - \hat{y}_n) + \frac{(1 - y_n)}{(1 - \hat{y}_n)} \hat{y}_n(1 - \hat{y}_n) \\ &= -y_n(1 - \hat{y}_n) + (1 - y_n)\hat{y}_n \\ &= -y_n + y_n\hat{y}_n + \hat{y}_n - y_n\hat{y}_n \\ &= \hat{y}_n - y_n \end{aligned}$$

# Natural combinations of nonlinearity and loss

Network Purpose	Activation of $k^{\text{th}}$ output unit	Loss for a single training data point	Gradient
Regression	$y_k = a_k$	$L = \frac{1}{2} \sum_{k=1}^K (\hat{y}_k - y_k)^2$	$\frac{\partial L}{\partial a_k} = \hat{y}_k - y_k$
Binary Classification	$y_k = \frac{1}{1 + \exp(-a_k)}$	$L = - \sum_{k=1}^K y_k \ln \hat{y}_k + (1 - y_k) \ln(1 - \hat{y}_k)$	$\frac{\partial L}{\partial a_k} = \hat{y}_k - y_k$
Multiclass Classification	$y_k = \frac{\exp(a_k)}{\sum_i \exp(a_i)}$	$L = - \sum_{k=1}^K y_k \ln \hat{y}_k$	$\frac{\partial L}{\partial a_k} = \hat{y}_k - y_k$



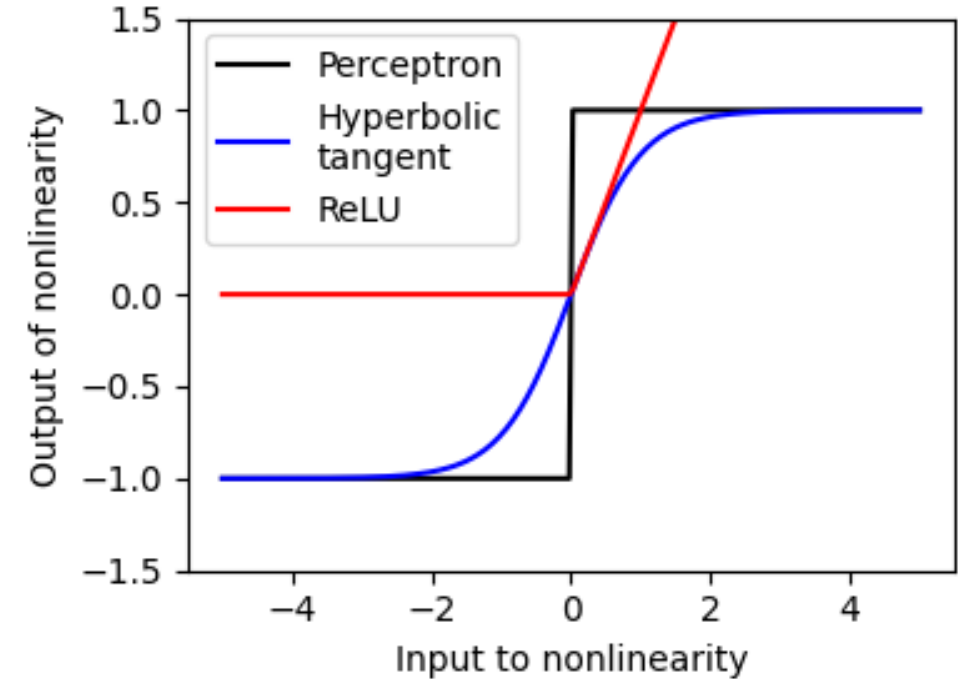
# Natural combinations of nonlinearity and loss

- This simple form for the derivative occurs whenever:
  - The target variable is assumed to have a conditional distribution from the exponential family
  - The loss for a training example is the negative log likelihood of the target
  - The activation function is chosen in a certain way (as the inverse of the canonical link function)
- For more detail see Bishop, *Pattern Recognition and Machine Learning*

**MULTI-LAYER NETWORKS ARE ALSO CALLED  
MULTI-LAYER PERCEPTRONS, BUT ARE NOT  
PERCEPTRONS**

# Terminology

- Artificial neural networks with at least two layers are often called multi-layer perceptrons
  - E.g., `sklearn.neural_network.MLPClassifier`
- However, such networks
  - Do not typically use the perceptron activation function (a step function)
  - Do not typically use the perceptron learning rule (they use gradient descent with different losses)
  - Can perform regression as well as classification
- Until about ten years ago, the most common activation function for hidden layers was the hyperbolic tangent, which is like a soft step function, so this made slightly more sense



# Summary

1. Multi-layer feedforward networks make a series of simple but unintuitive calculations
2. Multi-layer networks are composite functions
3. Multi-layer networks break the curse of dimensionality
4. A two-layer network can closely approximate any continuous function on a bounded domain
5. Networks with different weight matrices can perform identical mappings
6. Classification and regression employ different output and loss functions
7. Multi-layer networks are also called multi-layer perceptrons, but are not perceptrons