

# Memories

**Nachiket Kapre**

nachiket@uwaterloo.ca



UNIVERSITY OF  
**WATERLOO**

# Outline

- ▶ Need for memory
  - ▶ CPU model (instruction + data)
- ▶ Kinds of Memories
  - ▶ Flip-Flops – Register File – SRAM – DRAM
- ▶ RTL Syntax for memories
- ▶ Timing diagrams for memories


## Need for memory

- ▶ **History:** For computability, you need to store history of events, or store intermediate results. Think of Turing Machine + infinite tape.
  - ▶ Cannot rely on environment to provide storage

# Memory


- ▶ **Cost:** When implementing storage in hardware, you have a choice of implementations. Must know when to use which kind of memory.
  - ▶ FFs are fast parallel access but expensive. Single-cycle access to all FFs or registers is possible.
    - ▶ Register Files are collections of FFs with multiple read/write ports
  - ▶ SRAMs (Static RAMs) are compact but only access one element at a time. Single-cycle access is possible.
  - ▶ DRAMs (Dynamic RAMs) are cheapest, but require several cycles per access
- ▶ **Scheduling:** Timing behavior of different memories is different. Must think of memory properties when scheduling a datapath.

# Memory in CPUs



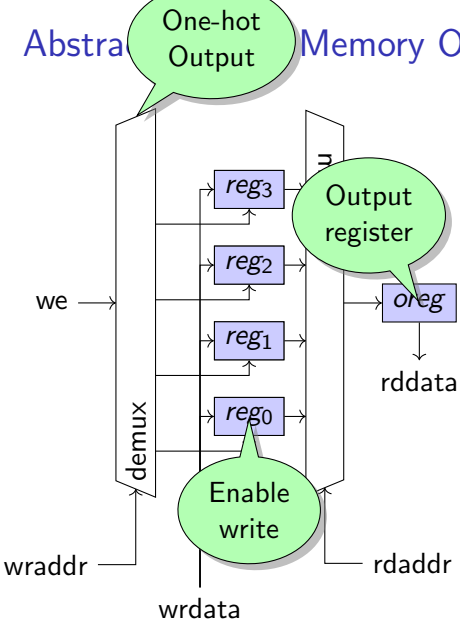
- ▶ Memory used to store instructions + data
- ▶ In hardware, instructions are just operations distributed in space + resource shared if required
- ▶ Variables on stack, dynamic allocation on heap → data memory
- ▶ Modern CPUs mix these two memory spaces, so we can write self-modifying code (downside: data can be made executable)
- ▶ For hardware design:
  - ▶ Focus on using dmem to store intermediate variables.
  - ▶ Implement imem for pre-compiled sequence of signals → microcode for datapath (state machine)

# Abstract View of Memory Operations





- ▶ Memory is a collection of storage locations (registers)
- ▶ **Mux** and **Demux** blocks arbitrate access the multiple memory locations
- ▶ Read address == select for the **Mux** block.
- ▶ Write address == select for the **Demux** block.
- ▶ Write enable = clock enable
- ▶ Reads are always happening, no enable needed

# Abstract Memory Operations



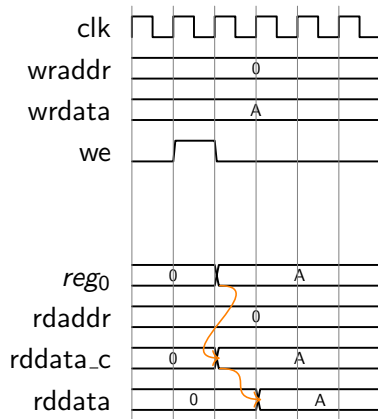

- ▶ Memory is a collection of storage locations (registers)
- ▶ **Mux** and **Demux** blocks arbitrate access the multiple memory locations
- ▶ Read address == select for the **Mux** block.
- ▶ Write address == select for the **Demux** block.
- ▶ Write enable = clock enable
- ▶ Reads are always happening, no enable needed

# Abstract View of Memory Operations





# Abstract View of Memory Operations



## Memory behavior

- ▶ Reads and writes can happen in parallel
  - ▶ FPGA RAMs allow both ports to be read/write ports
- ▶ Reads are happening **all** the time on the **rdaddr** location. No explicit read enable is needed
- ▶ Writes need explicit write enable which becomes a **ce** input to register (memory location)
- ▶ Write data gets stored into register in one cycle
- ▶ Read data available one cycle after address provided, or register state changed
  - ▶ If **rdaddr** stable, output data changes one cycle after **reg**
  - ▶ If **reg** values are stable, output data changes on cycle after **rdaddr**
- ▶ Thus, end-to-end delay write+read to+from same location = 2 cycles

# RTL Code for Memory Inference

```
module mem #(
    parameter [31:0] ADDRWIDTH=8,
    parameter [31:0] DATAWIDTH=32) (
    input wire clk,
    input wire rst,
    input wire [DATAWIDTH - 1:0] wrdata,
    output reg [DATAWIDTH - 1:0] rddata,
    input wire [ADDRWIDTH - 1:0] wraddr,
    input wire [ADDRWIDTH - 1:0] rdaddr,
    input wire we
);
//
reg [DATAWIDTH - 1:0] mem[2 ** ADDRWIDTH - 1:0];

integer i;
always @(posedge clk) begin
    if(rst) begin
        `ifndef SYNTHESIS //
            for (i=0; i <= 2**ADDRWIDTH - 1; i = i+1) begin
                mem[i] <= i;
            end
        `endif
        rddata <= 0;
    end else begin
        if(we) begin //
            mem[wraddr] <= wrdata;
        end
        rddata <= mem[rdaddr]; //
    end
end
end
endmodule
```

- ▶ Define an array with two parameters
  - ▶ number of data bits
  - ▶ number of address bits
- ▶ Resetting RAM only OK for simulations
- ▶ Clocked always block to manage **simultaneous** read and write operations
- ▶ Q: What happens if read+write to same addr in same cyc?

# RTL Code for Memory Inference

```
module mem #(
    parameter [31:0] ADDRWIDTH=8,
    parameter [31:0] DATAWIDTH=32) (
    input wire clk,
    input wire rst,
    input wire [0:7] wrdata,
    output reg [0:31] rddata,
    input wire [0:7] wraddr,
    input wire [0:7] rdaddr,
    input wire we,
    input wire rd);
// reg [DATAWIDTH-1:0] mem[0:255] = {0};

integer i;
always @(posedge clk) begin
    if(rst) begin
        `ifndef SYNTHESIS //
            for (i=0; i <= 255; i = i+1) begin
                mem[i] <= 0;
            end
        `endif
        rddata <= 0;
    end else begin
        if(we) begin //
            mem[wraddr] <= wrdata;
        end
        rddata <= mem[rdaddr]; //
    end
end
end
endmodule
```

2D

Initialization  
cannot be  
synthesized

Conditional  
Write

Unconditional  
Read

- ▶ Define an array with two parameters
  - ▶ number of data bits
  - ▶ number of address bits
- ▶ Resetting RAM only OK for simulations
- ▶ Clocked always block to manage **simultaneous** read and write operations
- ▶ Q: What happens if read+write to same addr in same cyc?

## Using Memories

- ▶ On-chip memories can be used to store intermediate data, inputs, and outputs
- ▶ To access a memory, you need to supply read/write address + write control → state machine
- ▶ Data to/from memories typically goes to user circuit → resource-shared or fully-pipelined datapath
- ▶ Address pattern + controls must match the expected data consumption rate of the datapath
  - ▶ A fully-pipelined circuit can consume + produce a set of inputs+outputs per cycle → state machine is simply a counter supplying address to memory
  - ▶ A resource-shared datapath consumes inputs once every THROUGHPUT number of cycles → need explicit state machine

# Using Memories with poly datapath (Fully-Pipelined)

- Compute the following function in hardware

```
#define N 1024


volatile int x[N]={1,2,3,4};
volatile int y[N]={0,0,0,0};

void poly(int a, int b, int c)
{
    for(int i=0;i<N;i++) {
        y[i]=a*x[i]*x[i]+b*x[i]+c;
    }
}


int main() {
    poly(1,2,3);
}
```

- Assume a, b, and c are inputs
- Store x and y into memory blocks
- Need a state-machine/counter to loop over items of x and y

## Using Memories with poly datapath (Fully-Pipelined)




## Using Memories with poly datapath (Fully-Pipelined)






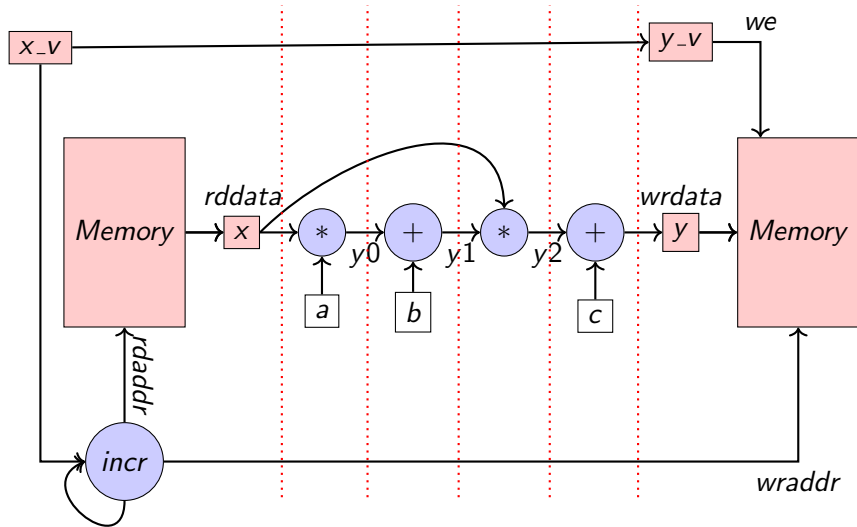
## Using Memories with poly datapath (Fully-Pipelined)



## Using Memories with poly datapath (Fully-Pipelined)



## Using Memories with poly datapath (Fully-Pipelined)



# Verilog Code for poly + Memory + Fully-Pipelined

```
poly_dp #(  
    .DATAWIDTH(DATAWIDTH))  
dp_inst (  
    .clk(clk),  
    .rst(rst),  
    .a(a),  
    .b(b),  
    .c(c),  
    .x(x),  
    .y(y));
```

```
poly_stmc #(  
    .ADDRWIDTH(ADDRWIDTH))  
stmc_inst(  
    .clk(clk),  
    .rst(rst),  
    .valid(x_v),  
    .rdaddr(rdaddr),  
    .wraddr(wraddr),  
    .we(we));
```

```
mem #(  
    .ADDRWIDTH(ADDRWIDTH),  
    .DATAWIDTH(DATAWIDTH))  
x_mem_inst(  
    .clk(clk),  
    .rst(rst),  
    .rddata(x),  
    .wrdata({DATAWIDTH{1'b0}}),  
    .wraddr({ADDRWIDTH{1'b0}}),  
    .rdaddr(rdaddr),  
    .we(1'b 0));
```

```
mem #(  
    .ADDRWIDTH(ADDRWIDTH),  
    .DATAWIDTH(3*DATAWIDTH))  
y_mem_inst(  
    .clk(clk),  
    .rst(rst),  
    .wrdata(y),  
    .wraddr(wraddr),  
    .rdaddr({ADDRWIDTH{1'b0}}),  
    .rddata(),  
    .we(we));
```

```
assign y_v = we;
```

# Verilog Code for State Machine

```
if(rst) begin
    valid_r <= 1'b 0; valid_r1 <= 1'b 0;
    valid_r2 <= 1'b 0; valid_r3 <= 1'b 0;
    valid_r4 <= 1'b 0;
    addr_r <= {ADDRWIDTH{1'b0}};
    addr_r1 <= {ADDRWIDTH{1'b0}};
    addr_r2 <= {ADDRWIDTH{1'b0}};
    addr_r3 <= {ADDRWIDTH{1'b0}};
    addr_r4 <= {ADDRWIDTH{1'b0}};
    addr_r5 <= {ADDRWIDTH{1'b0}};
end else begin
    valid_r <= valid;
    valid_r1 <= valid_r;
    valid_r2 <= valid_r1; //
    valid_r3 <= valid_r2;
    valid_r4 <= valid_r3;
    if(valid) begin //
        addr_r <= addr_r + 1;
    end
    addr_r1 <= addr_r;
    addr_r2 <= addr_r1;
    addr_r3 <= addr_r2;
    addr_r4 <= addr_r3; //
    addr_r5 <= addr_r4;
end
end

assign rdaddr = addr_r;
//
assign wraddr = addr_r5;
assign we = valid_r4;
```

- ▶ State machine generates addresses + write enable
- ▶ valid is input from external world
- ▶ rdaddr generated from a simple incrementer
  - ▶ incrementer controlled by valid
- ▶ wraddr and we are simply delayed versions of rdaddr and valid
  - ▶ Handled exactly like bubbles
- ▶ Notice that wraddr and we are offset by a cycle

# Verilog Code for State Machine

```
if(rst) begin
    valid_r <= 1'b 0; valid_r1 <= 1'b 0;
    valid_r2 <= 1'b 0; valid_r3 <= 1'b 0;
    valid_r4 <= 1'b 0;
    addr_r <= {ADDRWIDTH{1'b0}};
    addr_r1 <= {ADDRWIDTH{1'b0}};
    addr_r2 <= {ADDRWIDTH{1'b0}};
    addr_r3 <= {ADDRWIDTH{1'b0}};
    addr_r4 <= {ADDRWIDTH{1'b0}};
    addr_r5 <= {ADDRWIDTH{1'b0}};
end else begin
    valid_r <= valid;
    valid_r1 <= valid_r;
    valid_r2 <= valid_r1; //
    valid_r3 <= valid_r2;
    valid_r4 <= valid_r3;
    if(valid) begin //
        addr_r <= addr_r + 1;
    end
    addr_r1 <= addr_r;
    addr_r2 <= addr_r1;
    addr_r3 <= addr_r2;
    addr_r4 <= addr_r3; //
    addr_r5 <= addr_r4;
end
end

assign rdaddr = addr_r;
//
assign wraddr = addr_r5;
assign we = valid_r4;
```

valid  
pipeline

rdaddr  
generation

wraddr  
pipeline

- ▶ State machine generates addresses + write enable
- ▶ valid is input from external world
- ▶ rdaddr generated from a simple incrementer
  - ▶ incrementer controlled by valid
- ▶ wraddr and we are simply delayed versions of rdaddr and valid
  - ▶ Handled exactly like bubbles
- ▶ Notice that wraddr and we are offset by a cycle

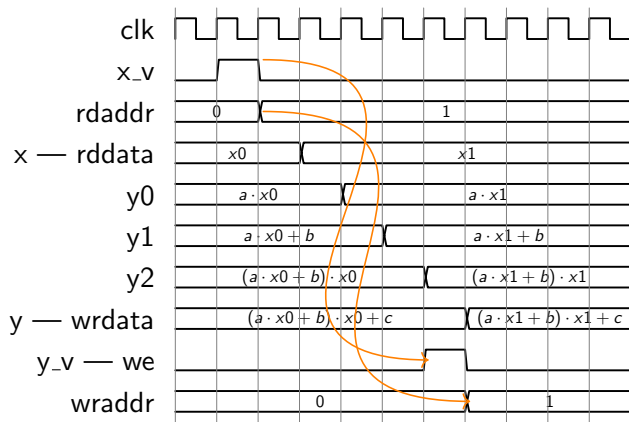
# Verilog Code for State Machine

```
if(rst) begin
    valid_r <= 1'b 0; valid_r1 <= 1'b 0;
    valid_r2 <= 1'b 0; valid_r3 <= 1'b 0;
    valid_r4 <= 1'b 0;
    addr_r <= {ADDRWIDTH{1'b0}};
    addr_r1 <= {ADDRWIDTH{1'b0}};
    addr_r2 <= {ADDRWIDTH{1'b0}};
    addr_r3 <= {ADDRWIDTH{1'b0}};
    addr_r4 <= {ADDRWIDTH{1'b0}};
    addr_r5 <= {ADDRWIDTH{1'b0}};
end else begin
    valid_r <= valid;
    valid_r1 <= valid_r;
    valid_r2 <= valid_r1; //
    valid_r3 <= valid_r2;
    valid_r4 <= valid_r3;
    if(valid) begin //
        addr_r <= addr_r + 1;
    end
    addr_r1 <= addr_r;
    addr_r2 <= addr_r1;
    addr_r3 <= addr_r2;
    addr_r4 <= addr_r3; //
end
//
assign rdaddr = addr_r;
//
assign wraddr = addr_r5;
assign we = valid_r4;
```

One cycle  
offset

- ▶ State machine generates addresses + write enable
- ▶ valid is input from external world
- ▶ rdaddr generated from a simple incrementer
  - ▶ incrementer controlled by valid
- ▶ wraddr and we are simply delayed versions of rdaddr and valid
  - ▶ Handled exactly like bubbles
- ▶ Notice that wraddr and we are offset by a cycle

## Timing Diagram of poly datapath





# Sharing Memories

- ▶ In the general case, consider memory ports as resources
- ▶ Track read/write operations on each memory port and disallow multiple reads/writes to the memory
  - ▶ Multi-ported memories are possible, but they are expensive, and should be used in high-performance designs only
- ▶ Generate address + write enables based on table

# Using Shared Memories with poly datapath

- ▶ Compute the following function in hardware

```
#define N 1024
```

```
volatile int a[N]={1,2,3,4};
```

```
volatile int b[N]={1,2,3,4};
```

```
volatile int c[N]={1,2,3,4};
```

```
volatile int x[N]={1,2,3,4};
```

```
volatile int y[N]={0,0,0,0};
```

```
void poly()
```

```
{
```

```
  for(int i=0;i<N;i++) {
```

```
    y[i]=a[i]*x[i]*x[i]+b[i]*x[i]+c[i];
```

```
  }
```

```
}
```


```
int main() {
```

```
  poly();
```

```
}
```

- ▶ Now, a, b, c and x are all inputs per iteration
- ▶ Cannot afford distinct memories for each array
- ▶ How do we resource-share the memory port across these arrays?
- ▶ Need a state-machine/counter to manage things properly

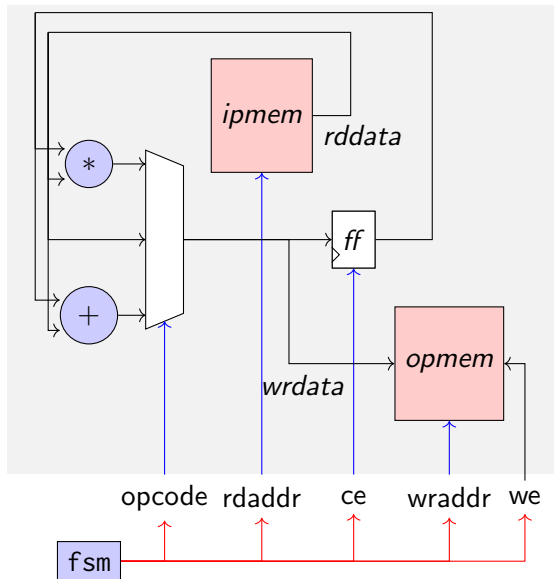
## Using Shared Memories with poly (Fully-Pipelined)



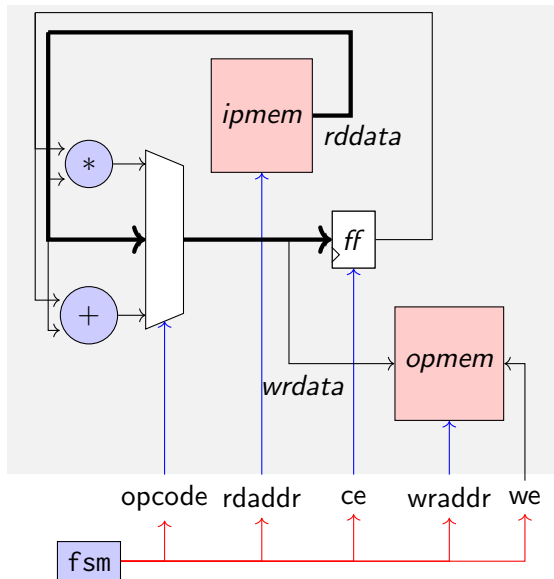
# Limits of using Shared Memories with Fully-Pipelined Datapaths

- ▶ Shared memories require multiple cycles to fetch/store inputs, outputs, or temporary variables in the computation
- ▶ A fully-pipelined/fully-spatial circuit would waste resources while accessing data from memory
- ▶ Resources will stay idle waiting for memory operations
- ▶ Engineering optimization → find system bottlenecks, reallocate resources to target bottleneck
- ▶ Can we reduce resource cost? → resource-shared datapath!

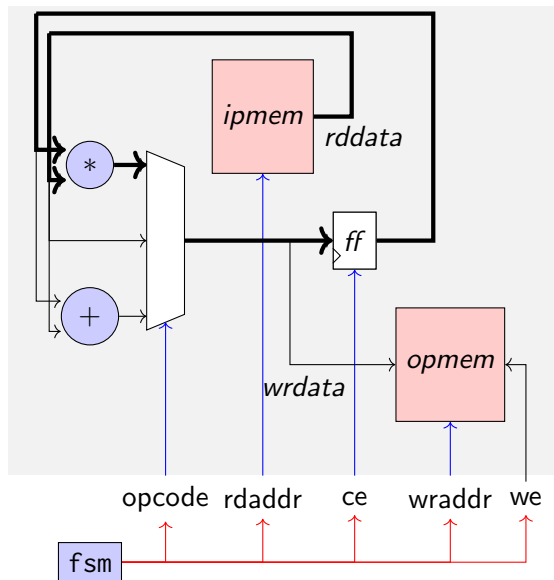
## Using Shared Memories with poly (Resource-Shared)



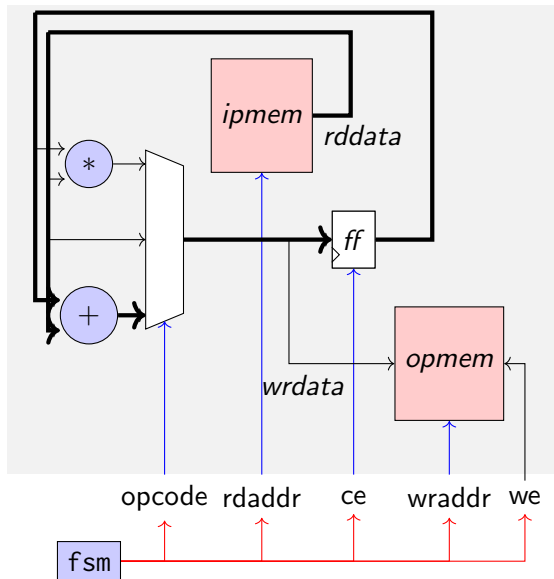
## Using Shared Memories with poly (Resource-Shared)



## Using Shared Memories with poly (Resource-Shared)




## Using Shared Memories with poly (Resource-Shared)

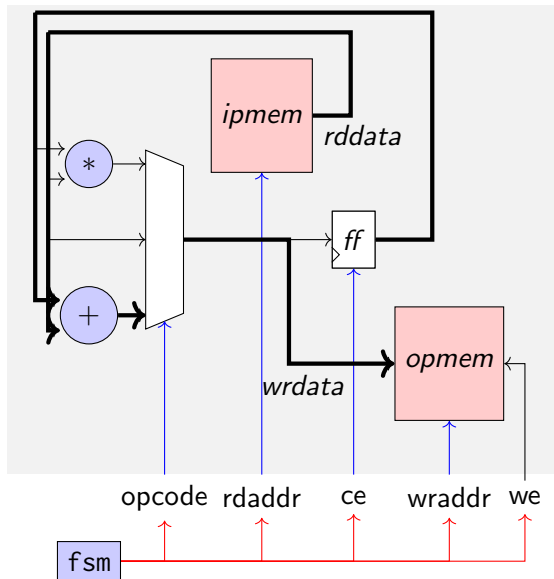




## Using Shared Memories with poly (Resource-Shared)



## Using Shared Memories with poly (Resource-Shared)




## Revisiting Tables

Cycle	Operators			
	$add_0$	$mult_0$	$load_0$	$store_0$
0	—	—	$a$	—
1	—	$reg_0 \cdot x$	$x$	—
2	$reg_0 + b$	—	$b$	—
3	—	$reg_0 \cdot x$	$x$	—
4	$reg_0 + c$	—	$c$	$y = add_0$


- ▶ Pack all input arrays  $x$ ,  $a$ ,  $b$ , and  $c$  into  $ipmem$ . Output  $y$  array stored in  $opmem$ .
- ▶ Modify schedule table to load  $a$  input in first cycle
- ▶ Modify datapath to provide a third “load” input to mux
- ▶ Note that no multiplexers required as input to operators  $\rightarrow$  all inputs in same memory
  - ▶ If multiple input memories used, muxes need to be put back

# Memory Layout

Interleaved



Block



## Address Generation (Interleaved)

Cycle	Memory Interface		
	<i>rdaddr</i>	<i>wraddr</i>	<i>we</i>
0	0 ( <i>a0</i> )	—	0
1	1 ( <i>x0</i> )	—	0
2	2 ( <i>b0</i> )	—	0
3	1 ( <i>x0</i> )	—	0
4	3 ( <i>c0</i> )	0 ( <i>y0</i> )	1

- ▶ Mux selection replaced by address generation
- ▶ Address to memory   muxsel to a set of memory location
- ▶ Data layout is an important concern
  - ▶ Interleaved *a0,x0,b0,c0,a1,x1,...*
  - ▶ Block *a0,a1,...,x0,x1,...*
  - ▶ Interleaved order simplifies address generation logic for this case
- ▶ Interleaving produces a simple state machine for generating read/write addresses.

# C-like code equivalence

```
#define N 1024

volatile int abcx[4*N]={1,2,3,4};
volatile int y[N]={0,0,0,0};

void poly()
{
    for(int i=0;i<N;i++) {
        int a = abcx[4*i+0];
        int x = abcx[4*i+1];
        int b = abcx[4*i+2];
        int c = abcx[4*i+3];
        y[i]=a*x*x+b*x+c;
    }
}

int main() {
    poly();
}
```

- ▶ Combine a, b, c, and x arrays into a single array abcx → sort-of mimic imem in hardware
- ▶ Assume interleaved packing a0,x0,b0,c0,a1,x1,b1,c1,...
- ▶ Must compute index into abcx as a function of loop index  $i$

# State Machine for Resource-Shared FSM

```
if(rst) begin
    counter <= {3{1'b0}};
    i <= {ADDRWIDTH{1'b0}};
end else begin
    counter <= counter + 1;
    if((counter == (TABLE_SIZE - 1))) begin
        counter <= {ADDRWIDTH{1'b0}};
        i <= i + 1;
    end
end
end
```

```
case(counter)
0 : begin
    rdaddr <= TABLE_SIZE * i + 0;
    wraddr <= {ADDRWIDTH{1'b0}};
    we <= 1'b 0;
end
1 : begin
    rdaddr <= TABLE_SIZE * i + 1;
    wraddr <= {ADDRWIDTH{1'b0}};
    we <= 1'b 0;
end
2 : begin
    rdaddr <= TABLE_SIZE * i + 2;
    wraddr <= {ADDRWIDTH{1'b0}};
    we <= 1'b 0;
end
3 : begin
    rdaddr <= TABLE_SIZE * i + 1;
    wraddr <= {ADDRWIDTH{1'b0}};
    we <= 1'b 0;
end
4 : begin
    rdaddr <= TABLE_SIZE * i + 3;
    wraddr <= i;
    we <= 1'b 1;
end
default : begin
    rdaddr <= TABLE_SIZE * i + 0;
    wraddr <= {ADDRWIDTH{1'b0}};
    we <= 1'b 0;
end
endcase
```

## Benefits of Resource Shared Datapath + Shared Memories

- ▶ Carefully balance resources against memory access bottleneck
- ▶ Isolated from rest of chip through ipmem and opmem structures
- ▶ Input and datapath muxes vanish → folded into the memory address decoders (internal to RAM structures)
- ▶ Now, we must create scheduling tables for memory read/write ports just like we did for operators and registers



# Wrapup

- ▶ Memories are efficient ways to store large number of registers
- ▶ Various kinds of memories on offer – multi-ported register files, SRAMs, DRAMs
- ▶ Full-throughput designs use exclusive memories for all arrays, no conflicts
- ▶ Resource-shared memories better match for resource-shared datapaths → area-time tradeoff